

WLE.rmd

David

Monday, August 18, 2014

Examination of Weight Lifting Exercises Dataset

Introduction

Six young healthy participants were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in five different fashions: exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E).

Their movements were monitored using four detectors strapped to their forearm, upperarm, waist and to the dumbbell itself. For feature extraction a sliding window approach was used with different lengths from 0.5 second to 2.5 seconds, with 0.5 second overlap. In each step of the sliding window approach the features on the Euler angles (roll, pitch and yaw), as well as the raw accelerometer, gyroscope and magnetometer readings were calculated.

The training data for this project are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

The data for this report comes from this source: <http://groupware.les.inf.puc-rio.br/har>.

Method

The algorithm selected to predict the quality of execution variable (classe) is random forest. This uses a decision tree approach with bootstrapping to generate a “forest” of potential trees. New trees are developed by restricting the number of variables that can be randomly selected at each node. This increases the variance between the trees improving the prediction accuracy.

The random forest methodology was chosen because: -It is unexcelled in accuracy among current algorithms for class prediction. -It runs efficiently on large data bases. -It can handle thousands of input variables without variable deletion. -It gives estimates of what variables are important in the classification.

Download and read files

The files are downloaded and read. The training file is named “training” and the testing file is named “validation”

```
if (!file.exists("data")) {  
  dir.create("data")  
}  
fileUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"  
fileUrltest <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"  
  
download.file(fileUrl, destfile = "./data/training.csv")  
download.file(fileUrltest, destfile = "./data/testing.csv")  
  
#Read the relevant tables.
```

```
training<- read.csv( "./data/training.csv")
validation <- read.csv( "./data/testing.csv")
```

Preprocess files

Some rows represent a summary of previous observations (the 2.5sec window as opposed to the 0.5). As the test data file does not provide this data our training and validation data files need to exclude them as well.

The measurements described in the introduction represent four measurement devices providing four measures each in three dimensions. This provides 48 predictors plus the predicted variable (classe). The next step of the preprocessing is to reduce the files to these 49 rows. This requires the removal of all extraneous variables (names, time, etc) which might give invalid correlations.

No further preprocessing was undertaken as the random forest method is not sensitive to issues of skewness etc..

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.1.1
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
training <- training[training$new_window=="no",]
training <- training[,c(8:10,37:48,60:68,84:86,113:124,151:160)]
nzv <- nearZeroVar(training,saveMetrics=T)
dim(training)
```

```
## [1] 19216    49
```

```
validation <- validation[,c(8:10,37:48,60:68,84:86,113:124,151:160)]
```

Finally we checked that the remaining predictors are valid by ensuring that they do not have a near zero variance (none do, see appendix). This gave a data frame of 19216 rows and 49 columns.

Create a training and testing set

The random forest uses a bootstrap method to estimates out of bag error (oob). This is used to estimate the optimal mtry (number of variables available for selection at each node)and hence does not provide an independent test of the accuracy of the final model.

Therefore the training file needs to be split into a training file (on which the model is built) and a teating file which will allow for cross validation and the calculation of the out of sample error

```
set.seed(1334)
```

```
inTraining <- createDataPartition(y=training$classe,p=0.6,list=FALSE)
training <- training[inTraining,]
testing <- training[-inTraining,]
```

Build model

As already explained the random forest method was adopted. However because the default approach was too computationally demanding for my PC the “traincontrol” argument was used. This provided for a 4 fold cross validation.

```
modFit <- train(classe~.,data=training,method="rf",trControl = trainControl(method = "cv", number = 4, a
```

```
## Loading required package: randomForest
```

```
## Warning: package 'randomForest' was built under R version 3.1.1
```

```
## randomForest 4.6-10
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
## Warning: package 'e1071' was built under R version 3.1.1
```

```
## + Fold1: mtry= 2
```

```
## - Fold1: mtry= 2
```

```
## + Fold1: mtry=25
```

```
## - Fold1: mtry=25
```

```
## + Fold1: mtry=48
```

```
## - Fold1: mtry=48
```

```
## + Fold2: mtry= 2
```

```
## - Fold2: mtry= 2
```

```
## + Fold2: mtry=25
```

```
## - Fold2: mtry=25
```

```
## + Fold2: mtry=48
```

```
## - Fold2: mtry=48
```

```
## + Fold3: mtry= 2
```

```
## - Fold3: mtry= 2
```

```
## + Fold3: mtry=25
```

```
## - Fold3: mtry=25
```

```
## + Fold3: mtry=48
```

```
## - Fold3: mtry=48
```

```
## + Fold4: mtry= 2
```

```
## - Fold4: mtry= 2
```

```
## + Fold4: mtry=25
```

```
## - Fold4: mtry=25
```

```
## + Fold4: mtry=48
```

```
## - Fold4: mtry=48
```

```
## Aggregating results
```

```
## Selecting tuning parameters
```

```
## Fitting mtry = 25 on full training set
```

```
print(modFit)
```

```
## Random Forest
```

```
##
```

```
## 11532 samples
```

```
##    48 predictors
```

```
##      5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 8649, 8648, 8649, 8650
##
## Resampling results across tuning parameters:
##
##      mtry  Accuracy  Kappa  Accuracy SD  Kappa SD
##      2      1          1      9e-04          0.001
##     20      1          1      7e-04          9e-04
##     50      1          1      0.002          0.003
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 25.
```

The model indicated the optimal mtry was 25 and this adopted model gives an accuracy of 0.988.

Predict result on test data

To calculate out of sample error this model was applied to the testing data and its results compared to the actual classe values using the confusion matrix function.

```
pred <- predict(modFit,testing)
print(confusionMatrix(pred,testing$classe))
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A    B    C    D    E
##      A 1345    0    0    0    0
##      B    0  863    0    0    0
##      C    0    0  818    0    0
##      D    0    0    0  731    0
##      E    0    0    0    0  867
##
## Overall Statistics
##
##              Accuracy : 1
##              95% CI : (0.999, 1)
##      No Information Rate : 0.291
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 1
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.000    1.000    1.000    1.000    1.000
## Specificity          1.000    1.000    1.000    1.000    1.000
```

## Pos Pred Value	1.000	1.000	1.000	1.000	1.000
## Neg Pred Value	1.000	1.000	1.000	1.000	1.000
## Prevalence	0.291	0.187	0.177	0.158	0.188
## Detection Rate	0.291	0.187	0.177	0.158	0.188
## Detection Prevalence	0.291	0.187	0.177	0.158	0.188
## Balanced Accuracy	1.000	1.000	1.000	1.000	1.000

The confusion matrix predicts 0 out of sample error and therefore there is no opportunity to improve the model.

Predict results for 20 test cases

Finally the model is used to predict the classe variable for the 20 test cases provided. It correctly predicts all 20.

```
answers <- predict(modFit,validation)
print(answers)
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Appendix

```
print(nzv)
```

##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.086	6.83805	FALSE	FALSE
## pitch_belt	1.037	9.53893	FALSE	FALSE
## yaw_belt	1.047	10.10616	FALSE	FALSE
## gyros_belt_x	1.051	0.72336	FALSE	FALSE
## gyros_belt_y	1.149	0.35908	FALSE	FALSE
## gyros_belt_z	1.071	0.87948	FALSE	FALSE
## accel_belt_x	1.059	0.85346	FALSE	FALSE
## accel_belt_y	1.115	0.74417	FALSE	FALSE
## accel_belt_z	1.081	1.55600	FALSE	FALSE
## magnet_belt_x	1.089	1.70171	FALSE	FALSE
## magnet_belt_y	1.097	1.55079	FALSE	FALSE
## magnet_belt_z	1.019	2.37302	FALSE	FALSE
## roll_arm	51.154	13.75937	FALSE	FALSE
## pitch_arm	85.282	15.96066	FALSE	FALSE
## yaw_arm	32.282	14.89904	FALSE	FALSE
## gyros_arm_x	1.024	3.34617	FALSE	FALSE
## gyros_arm_y	1.451	1.95150	FALSE	FALSE
## gyros_arm_z	1.119	1.29059	FALSE	FALSE
## accel_arm_x	1.018	4.04351	FALSE	FALSE
## accel_arm_y	1.169	2.78414	FALSE	FALSE
## accel_arm_z	1.139	4.12157	FALSE	FALSE
## magnet_arm_x	1.012	6.96295	FALSE	FALSE
## magnet_arm_y	1.045	4.53268	FALSE	FALSE
## magnet_arm_z	1.028	6.57785	FALSE	FALSE

## roll_dumbbell	1.038	83.75312	FALSE	FALSE
## pitch_dumbbell	2.248	81.22398	FALSE	FALSE
## yaw_dumbbell	1.132	83.02456	FALSE	FALSE
## gyros_dumbbell_x	1.010	1.25416	FALSE	FALSE
## gyros_dumbbell_y	1.271	1.44151	FALSE	FALSE
## gyros_dumbbell_z	1.053	1.06682	FALSE	FALSE
## accel_dumbbell_x	1.006	2.21170	FALSE	FALSE
## accel_dumbbell_y	1.062	2.41986	FALSE	FALSE
## accel_dumbbell_z	1.150	2.12323	FALSE	FALSE
## magnet_dumbbell_x	1.094	5.84929	FALSE	FALSE
## magnet_dumbbell_y	1.189	4.38177	FALSE	FALSE
## magnet_dumbbell_z	1.027	3.51270	FALSE	FALSE
## roll_forearm	11.726	11.23543	FALSE	FALSE
## pitch_forearm	64.576	15.09679	FALSE	FALSE
## yaw_forearm	15.236	10.29871	FALSE	FALSE
## gyros_forearm_x	1.050	1.54559	FALSE	FALSE
## gyros_forearm_y	1.043	3.84055	FALSE	FALSE
## gyros_forearm_z	1.112	1.58201	FALSE	FALSE
## accel_forearm_x	1.143	4.13197	FALSE	FALSE
## accel_forearm_y	1.050	5.20920	FALSE	FALSE
## accel_forearm_z	1.019	3.01311	FALSE	FALSE
## magnet_forearm_x	1.013	7.92569	FALSE	FALSE
## magnet_forearm_y	1.256	9.72627	FALSE	FALSE
## magnet_forearm_z	1.018	8.75833	FALSE	FALSE
## classe	1.471	0.02602	FALSE	FALSE

References

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13) . Stuttgart, Germany: ACM SIGCHI, 2013